

CS-453 - Project

Software

Transactional Memory




Distributed Computing Laboratory

September 26 2023

In short:


- 30% of the grade of CC
- A single deadline
- Individual submission (YOUR code)
- But we suggest you work in small groups to help each other
- Reference document on Moodle
- Automated grader
- Project sessions are Q&A

Let me introduce you to Alice...

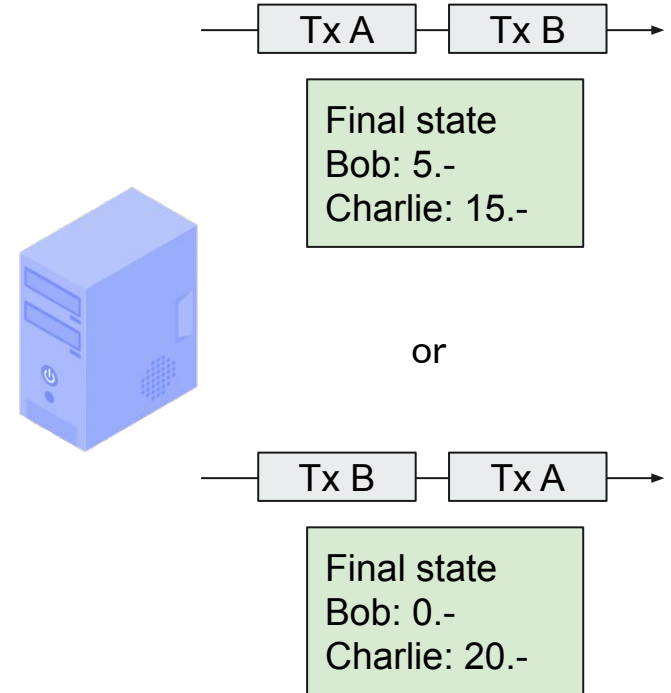





How Alice's Bank works


```
integer[] accounts;  
  
fn transfer(src, dst, amount) {  
    if (accounts[src] < amount) // Not enough funds  
        return; // => no transfer  
    accounts[dst] = accounts[dst] + amount;  
    accounts[src] = accounts[src] - amount;  
}
```





Executions



Scaling



Free lunch is over



Closer look



```
integer[] accounts;  
  
fn transfer(src, dst, amount) {  
    let v0 = accounts[src]; // 0th read  
    if (v0 < amount)  
        return;  
    let v1 = accounts[dst]; // 1st read  
    accounts[dst] = v1 + amount; // 1st write  
    let v2 = accounts[src]; // 2nd read  
    accounts[src] = v2 - amount; // 2nd write  
}
```

Concurrent transfers

Initial state

Bob: 20.-CHF
Charlie: 0.-CHF

Tx A
FROM Bob
TO Charlie
AMOUNT 20.-CHF

Read B == 20;
Read C == 0;
Write C = 20;



Final state

Bob: 25.-CHF
Charlie: 15.-CHF

Tx B
FROM Charlie
TO Bob
AMOUNT 5.-CHF

Read C == 20;
Read B == 20;

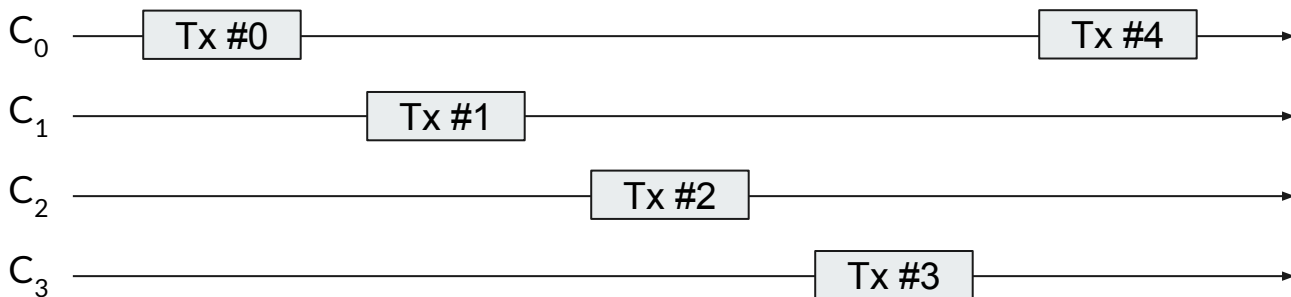
Read B == 20;
Write B = 0;

Write B == 25;
Read C == 20;
Write C = 15;

A quick fix


```
integer[] accounts;
```

```
fn safe_transfer(src, dst, amount) {  
    lock accounts { // Only one CORE can access accounts at a time  
        fn transfer(src, dst, amount);  
    }  
}
```



Coarse-grained locking is safe, but doesn't scale.

Parallel transfers



(Santa Bob and Santa Charlie live in a different universe and don't trade with Boring Bob and Boring Charlie.) ¹⁰

Alice's wish list



Alice would like an abstraction that:

- Is (almost) as simple to use as a coarse-grained lock,
- Allows non-conflicting *transactions* (transfers in our case) to run in parallel,
- Serializes conflicting transactions.

Well, that's exactly Transactional Memory!

Informally, a memory transaction is



- A sequence of actions (READs/WRITEs/ALLOCs/FREEs),
- That execute **atomically** (partially-executed transactions are not visible),
- And either **all** actions **commit** or **abort together** (all or nothing).

A Software Transactional Memory (STM) is a software framework that lets one write and run memory transactions.

Safe transactions using the STM

```
// Initialization of the transactional memory, which holds the accounts
tm = new TransactionalMemory(/*...*/);
integer[] accounts = tm.get_start();

fn transfer(src, dst, amount) {
    while (true) { try {
        let tx = tm.begin();
        if (tx.read(&accounts[src]) < amount) {
            tx.commit();
            break;
        }
        tx.write(&accounts[dst], tx.read(&accounts[dst]) + amount);
        tx.write(&accounts[src], tx.read(&accounts[src]) - amount);
        tx.commit();
        break;
    } except (RetryTransaction) { continue; } }
}
```

Implement a STM to help Alice's business grow!

9 functions to implement:

- `tm_create / tm_destroy` // constructor and destructor for TM.
- `tm_start` // returns an opaque pointer to the start of the TM.
- `tx_begin / tx_end` // starts/tries to commit a transaction.
- `tx_read` // reads a value targeted by an opaque pointer (in a transaction).
- `tx_write` // writes a value to an opaque pointer (in a transaction).
- `tx_alloc` // allocates a new memory buffer and returns an opaque pointer (in a transaction).
- `tx_free` // frees a previously allocated buffer (in a transaction).

Correctness of your implementation



Informally, your STM should make concurrent memory transactions appear as if they were executed serially, without concurrency.

In order to be correct, your implementation must have the 3 following properties:

- Snapshot isolation: After its start, a transaction cannot see (*via* reads/frees) modifications from other concurrent transactions (*via* writes/allocs/frees).
- Atomicity: All modifications from a transaction appear to take place at one indivisible point in time.
- Consistency: Committed transactions continue with the state left by the last committed one (according to the *linearization* order).

Resources



- *The project reference document (19-page pdf) on Moodle:*
 - Everything in this presentation,
 - Formal specs of all functions, properties, etc.
 - Possible implementation (with pseudocode) (without data structures),
 - Instructions to test and submit your work.
- GitHub repo: <https://github.com/LPD-EPFL/CS453-2022-project>
 - `include/` headers
 - `template/` where to work
 - `reference/` a correct but SLOW implementation that uses a global lock
 - `grading/` to grade your work
 - a python script to submit your work

Grading (1/2)

- 30% of the grade of Concurrent Algorithms.
- Correctness is a MUST: incorrect => no passing grade for the project (< 4)
- FORBIDDEN to have an implementation similar to the reference one (i.e., that uses a coarse lock).
- FORBIDDEN to never let non-conflicting transactions run in parallel.
- Your grade depends on the speedup vs the reference implementation.
- 16x slower than the reference version => no passing grade.
- Projects with memory leaks are capped at 5.

Speedup s	Grade g
$s < 1/16$	$g < 4$
$s \in [1/16, 1]$	$g = 4 + \frac{16s-1}{15}$
$s \in [1, 2.5]$	$g = 5 + \frac{2s-2}{3}$
$s > 2.5$	$g = 6$

Grading (2/2)



- Automated submission system with unlimited submissions allowed until the deadline (only your best submission is considered for grading).
- You can collaborate with other students (debug each other's code, share ideas, etc.) and take inspiration from existing STM libraries, but you should submit YOUR OWN implementation.
- Any attempt to CHEAT (e.g., tricking the grader, plagiarism) will be SEVERELY PUNISHED.

Important notes



- You can use C or C++. C++ won't give you much advantage: unless you prefer it, go with C.
- You can use any OS or compiler for your local development, but be aware I may not be able to help you.
- Only Ubuntu and Debian, with reasonably recent GCC/clang, are supported troubleshooting-wise.
- You are not allowed to use libraries that solve (non-elementary) concurrent programming problems.
- Check that you received an email from me with your UUID, tell me ASAP if you didn't!

Where to go



- Read *the* pdf.
 - Read it again.
 - Clone the repository.
 - Understand the reference (bad) implementation.
 - Implement the algorithm proposed in the pdf.
 - Test/grade locally until you're satisfied.
 - Submit your work: done. :)
-
- Few locs (~1000), but tricky ones.
 - Debugging concurrent code is VERY hard, START as SOON as possible.
-
- Seriously, for your own good, start now.

Questions?