



# **CS-453 - Project**

(An overview of)

# Concurrent Programming in C/C++

Distributed Computing Laboratory

September 26 2023

# Back to CS101



```
// a single thread

int a = 0;
int b = 0;
print(a, b);      // a = 0, b = 0

a = 1;
print(a, b);      // a = 1, b = 0

b = 1;
print(a, b);      // a = 1, b = 1
```

# What if we have two threads?


```
// Global var.      // Thread B

int a = 0;           int v = b; // read
int b = 0;           if (v==1) {
                    print(a, v);
// Thread A         // a = 1, v = 1?
                    // a = 1, v = 0?
a = 1; // write      // a = 0, v = 1?
b = 1; // write      // a = 0, v = 0?
                    }
}
```

According to common sense / "sequential consistency"	What will happen in C/C++	What could happen according to the C/C++ standards
		Format your disk


# Who are the culprits?

1) Your **C/C++ compiler** can **reorder** instructions if it **doesn't** have any **local side effects**.



2) According to the C/C++ standards, **accessing a variable** that is being written by another thread **without synchronization** (data race) is an **Undefined Behavior**, it can lead to absolutely anything.

3) Your **CPU**, depending on its consistency model, can execute **unrelated R/W** (& R/R, W/W) **out-of-order**.



But, when programming in C/C++, you shouldn't have to care about what your hardware promises, only the C/C++ standards.

# The main takeaway



C/C++ do **NOT** ensure (without extra care)

that reads/writes

are carried/observed

**in program order**

by **different threads**

You need to use synchronization primitives when sharing data across threads to restore sequential consistency.

# Example: let's build a concurrent counter



```
#include <pthread.h>
#include <assert.h>

static int counter = 0;

void* thread(void* null) {
    counter = counter + 1; // race condition
}
```

```
int main() {
    pthread_t handlers[2];
    for (int i = 0; i < 2; i++)
        pthread_create(&handlers[i], NULL, thread, NULL);
    for (int i = 0; i < 2; i++)
        int res = pthread_join(handlers[i], NULL);
    assert(counter == 2);
}
```

Let's try to fix this example by using synchronization primitives!

# Sync primitive #1: Locks/Mutexes

- A lock (or **Mutual Exclusion**) can only be held by one thread at a time.
- Use it to **prevent data races** on shared variables.
- It will **prevent reordering** via a **fence** and ensure **sequential consistency**.

```
#include <pthread.h>
```

```
#include <assert.h>
```

```
pthread_mutex_t mutex;
```

```
static int counter = 0;
```

```
void* thread(void* null) {
```

```
    pthread_mutex_lock(&mutex);
```

```
    counter = counter + 1;
```

```
    pthread_mutex_unlock(&mutex);
```

```
}
```

```
int main() {
```

```
    pthread_mutex_init(&mutex, NULL);
```

```
    pthread_t handlers[2];
```

```
    for (int i = 0; i < 2; i++)
```

```
        pthread_create(&handlers[i], NULL, thread, NULL);
```

```
    for (int i = 0; i < 2; i++)
```

```
        int res = pthread_join(handlers[i], NULL);
```

```
    assert(counter == 2);
```

```
    pthread_mutex_destroy(&mutex);
```

```
}
```

## Sync primitive #2: Atomic variables (1/2)

- An atomic variable can safely be accessed concurrently from multiple threads (no data races)
- They offer **atomic operations** (i.e., no other thread can observe partially-completed ops):
  - Read (atomic\_load) / Write (atomic\_store)
  - Increment (atomic\_fetch\_add) / Compare and Swap (atomic\_compare\_exchange\_strong)
- (By default,) They prevent reorderings and offer **sequential consistency**.

```
#include <pthread.h>
#include <assert.h>
#include <stdatomic.h>

static atomic_int counter = 0;

void* thread(void* null) {
    atomic_fetch_add(&counter, 1);
}
void* bad_thread(void* null) {
    counter = counter + 1; // 2 ATOMIC OPERATIONS (LOAD and STORE) INSTEAD OF 1
}
```

```
int main() {
    pthread_t handlers[2];
    for (int i = 0; i < 2; i++)
        pthread_create(&handlers[i], NULL, thread, NULL);
    for (int i = 0; i < 2; i++)
        int res = pthread_join(handlers[i], NULL);
    assert(counter == 2);
}
```



## Sync primitive #2: Atomic variables (2/2)

Atomic variables can be used to implement locks using the “Compare and Swap” operation

```
#include <stdatomic.h>
```

```
#define UNLOCKED 0
```

```
#define LOCKED 1
```

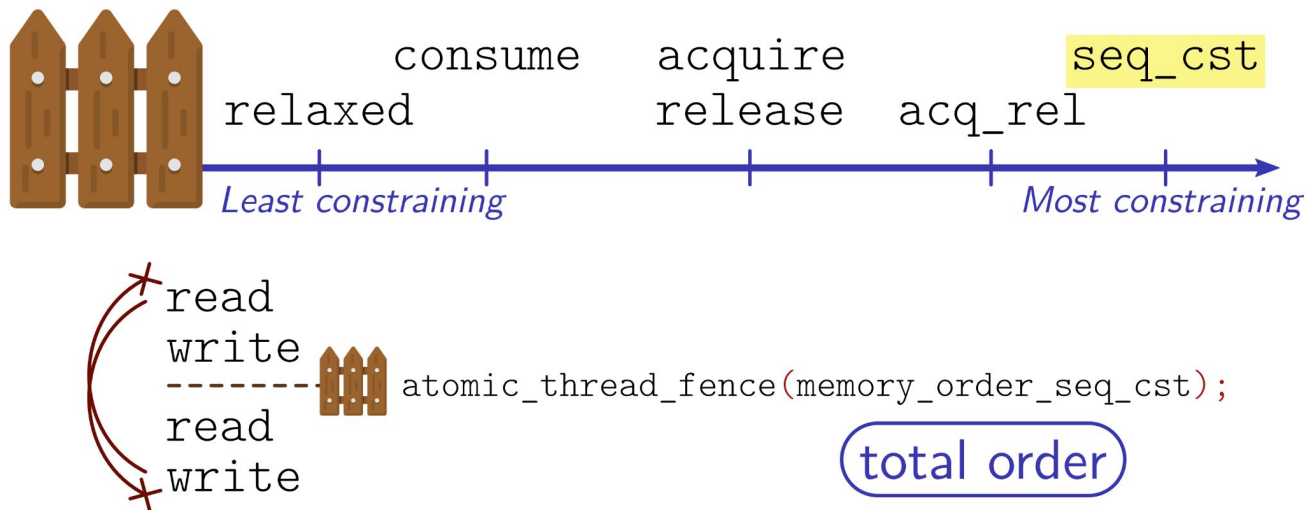
```
struct lock {  
    atomic_bool state;  
};
```

```
void init_lock(struct lock* lock) {  
    lock->state = UNLOCKED;  
}
```

```
void take_lock(struct lock* lock) {  
    while (true) {  
        bool expected = UNLOCKED;  
        atomic_compare_exchange_strong(  
            &lock->state, &expected, LOCKED);  
        if (expected == UNLOCKED) break;  
    }  
}
```

```
void release_lock(struct lock* lock) {  
    lock->state = UNLOCKED;  
}
```

# Sequential Consistency is a strict ordering



- Sequential Consistency prevents all reordering and can become a bottleneck.
- You can make your program more efficient by allowing some reordering.
- But it gets tricky to reason about, and you probably won't need it for this class. :)
- [https://en.cppreference.com/w/c/atomic/memory\\_order](https://en.cppreference.com/w/c/atomic/memory_order)